

Does a present bias influence exploratory choice?

Alexander S. Rich (asr443@nyu.edu)

Todd M. Gureckis (todd.gureckis@nyu.edu)

New York University, Department of Psychology, 6 Washington Place, New York, NY 10003 USA

Abstract

Balancing exploration and exploitation is difficult, and across a wide variety of situations under-exploration of uncertain alternatives appears prevalent. We propose that one possible cause of under-exploration is present bias, whereby immediate rewards (like those gained from exploitation) loom larger than future rewards (like those gained from exploration). This possible cause of under-exploration is not addressed by past studies, in which choices generally yield token rewards that are converted to money at the end of the experiment, removing the inter-temporal aspect of the decision-making process. In this paper, we develop an exploratory choice task with immediately-consumed rewards. In Experiment 1, we show that people understand the task and respond as expected to standard decision variables. In Experiment 2, we introduce a condition in which there is a temporal delay between choices and outcomes, and test whether this increases exploration by reducing present bias. We find no evidence of an effect of delay on exploration, and discuss possible interpretations of this result.

Keywords time preference, present bias, exploratory choice

Balancing exploitation of known options with exploration of uncertain options presents a difficult cognitive and motivational challenge. The problem of how to trade off between these two conflicting goals has created whole fields of study (Mehlhorn et al., 2015; Sutton & Barto, 1998), and in many cases achieving an optimal balance is computationally intractable (Guez et al., 2013). Given these facts, one might expect people to often deviate from optimal exploratory choice in both directions, perhaps exploring too much in some situations but too little in others. Notably, though, while over-exploration does occur, a bias towards under-exploration in particular seems to be observed across a wide range of domains.

In learned helplessness, for example, an organism experiences the absence of control over the environment, learns that the environment is uncontrollable, and thus ceases to take actions that might allow it to discover that it can in fact exert control. While the cognitive appraisal of experienced events affects the development of learned helplessness (Abramson et al., 1978), patterns of exploration clearly play a role as well (Huys & Dayan, 2009; Teodorescu & Erev, 2014a).

Under-exploration also seems to occur in the development of complex skills, such as flying a plane or playing a sport (Gopher et al., 1989). In these settings, an “emphasis change” training method that encourages people to continually explore the performance space leads to greater performance gains than unguided practice or more complex training methods. Without this intervention, people often enter a “local maximum” in which exploration decreases and performance plateaus (Yechiam et al., 2001).

In many other areas under-exploration is less clearly established, but is suspected to play a role in maladaptive behavior.

Insufficient exploratory interaction with outgroups may be one cause of stereotypes and prejudice (Denrell, 2005), and interventions that increase inter-group contact reduce stereotypes (Shook & Fazio, 2008). The crowding out of exploration by exploitation is a concern in organizational behavior as well (March, 1991; Levinthal & March, 1993), prompting research into organizational structures that may preserve exploration (Fang et al., 2010).

Under- and over-exploration in the lab

Given the seeming pervasiveness of under-exploration in real-world decision making, one might expect a clear pattern of under-exploration in lab experiments. Instead, lab results are mixed, with people sometimes under-exploring, sometimes over-exploring, and sometimes exploring close to an optimal amount. To take two illustrative examples, Zwick et al. (2003) found that in a sequential search task people under-searched when there were no information costs but over-searched when there were information costs, and Teodorescu & Erev (2014b) found that people explored unknown alternatives too often or not often enough depending on whether rare outcomes were positive or negative. Similar results have been obtained within and across a variety of other studies and paradigms (Tversky & Edwards, 1966; Hertwig et al., 2004; Navarro et al., 2016; Juni et al., 2016; Sang et al., 2011).

These experimental studies raise the question of why under-exploration appears more widespread than over-exploration in the field, but not in the lab. One possibility is that both forms of deviation from optimality are in fact prevalent, though perhaps in different settings, and that the seemingly general bias toward under-exploration is illusory. In the current paper, we propose an alternative hypothesis: that lab tests of exploratory choice have been missing a key component of real-world exploratory choice, namely, the distribution of choices and outcomes over time, and that because of people’s bias towards immediate rewards, this spreading of choices over time may account for the tendency to under-explore.

Exploration and myopia

Past theoretical work has proposed that under-exploration may be the result of myopia: the underweighting of temporally distant rewards relative to close ones (March, 1991; Levinthal & March, 1993). In expectation, exploitation yields the greatest reward for the present choice; exploration yields lower expected reward for the present choice, but raises the expected reward for all future choices due to the possibility of discovering a better option. Thus, if future choices and rewards are weighted less than immediate ones, the value of

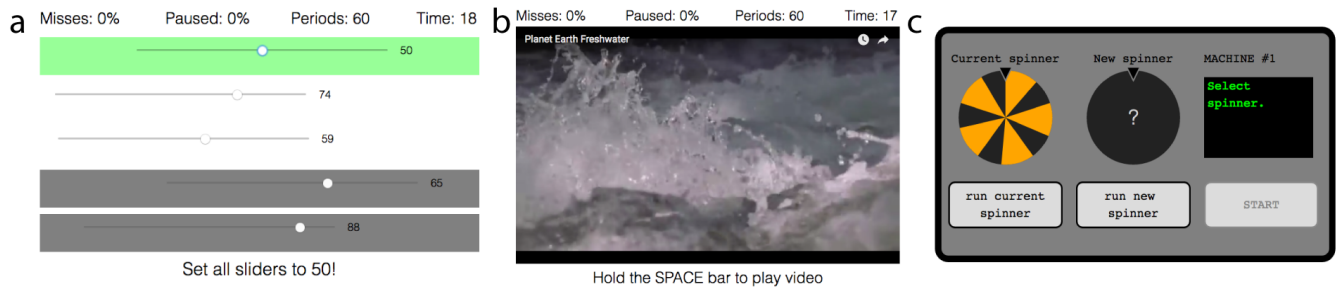


Figure 1: *a*: An example of the slider task. Participants had to move all sliders to “50” in 30 seconds. *b*: the video-watching task. Participants had to hold the space bar to watch their chosen video. *c*: The decision-making task. Participants had to choose to run the machine with the current spinner or try a new spinner. If their chosen spinner landed on a gold wedge, they performed the video-watching task instead of the slider task.

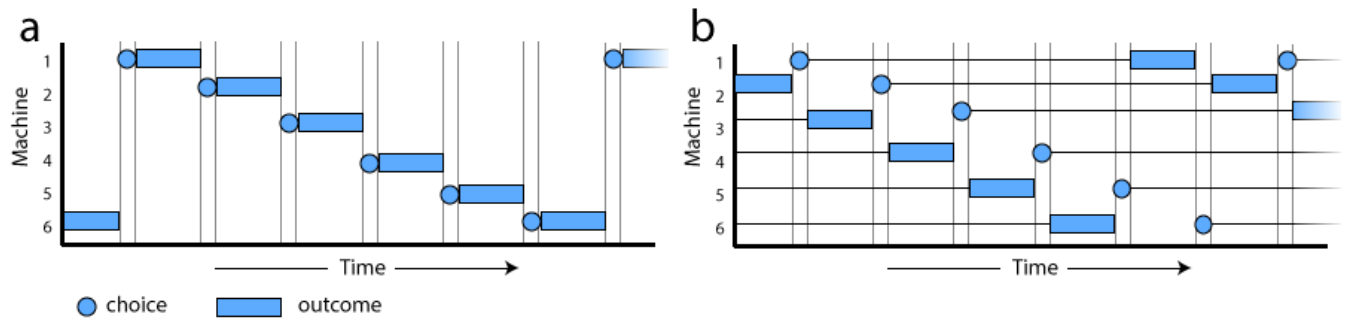


Figure 2: A schematic illustration of the experiment timeline. In all experiments and conditions, six 30-second work periods passed between consecutive choices with a given machine. *a*: in Experiment 1 and the immediate condition of Experiment 2, the outcome for a choice was revealed and occurred immediately after the choice was made. *b*: in the delayed condition of Experiment 2 (right), the outcome of a choice was revealed and occurred after a delay of four work periods.

exploration is diminished.

There are two distinct possible forms of myopia (Bartels & Urminsky, 2015). The first is that the decision-maker is unaware of or fails to consider the future. This form of myopia is likely one cause of under-exploration, particularly in complex environments. However, researchers in exploratory choice have begun to address this form of myopia, and have found that people change their patterns of exploration based on their expectations about future choices, showing that, at least to some degree they do consider the future (Wilson et al., 2014; Rich & Gureckis, 2017).

The second form of myopia is that the decision-maker is aware of the future but simply doesn't care as much about it. A wealth of studies have shown that this form of myopia is an important factor in human decision making, and that people are biased towards the present. For example, people will often prefer a larger, later monetary reward to a smaller, sooner reward when both rewards are in the future, but will switch their preference when the time until both rewards is reduced so that the sooner reward is immediate (Kirby & Herrnstein, 1995). With monetary rewards, the delay or speed-up required to observe preference reversals is usually several days. With non-monetary rewards, such as the cessation of an annoying noise (Solnick et al., 1980), watching a video when bored (Navarick, 1998), or drinking soda when thirsty (Brown et al., 2009), a bias towards immediate rewards has been observed on the scale of minutes or even seconds.

Present bias has the clear potential to cause people to explore too little, by increasing the weight of the greater expected immediate gains of exploiting relative to the greater expected delayed gains of exploring. However, there has been no experimental research examining how people's bias towards the present affects exploratory choice. In part, this is due to limits in experiment design. In most exploratory choice experiments, each choice and outcome takes fewer than 10 seconds to complete, making it unlikely for a large bias to develop between the current trial and future ones. More importantly, regardless of the length of each trial, rewards are generally not consumed on a trial-by-trial basis. For example, a participant might receive “points” on each trial, which are converted to money when the experiment ends. In this sort of paradigm, there is no reason to expect a bias towards the present choice's rewards over future rewards, as all rewards are in fact delivered at the same time.

In the present study, we describe our ongoing efforts to study the effect of present bias on exploratory choice. We develop a novel exploratory choice task using immediately-consumed rewards and present two experiments. In the first, we vary the expected payoff of exploration across conditions and test whether this affects participant choice. The results of this experiment show that participants understand the task and respond as predicted to well-studied variables. In the second, we introduce a delayed-outcome condition. If people are present biased, we predict this to increase exploration by de-

creasing the weight of the current choice's outcomes relative to future choice's outcomes. However, we find no strong evidence that delaying outcomes affects the rate of exploration. While this may support the conclusion that exploratory choice is not susceptible to present bias, we also discuss other possible explanations of this null result and plans for future experiments.

Experiment 1

In our experimental paradigm, participants made a series of explore–exploit decisions, where each decision determined whether the following 30-second consumption task would be unpleasant (moving sliders) or pleasant (watching a YouTube video). In Experiment 1, we sought to confirm that participants preferred the video task to the slider task and would seek to increase video-watching time, and that people were sensitive to explicit, exogenous changes to the value of exploration. The experiment consisted of two conditions—in the high-explore condition, the average payoff of exploration was high, and in the low-explore condition, it was low. We predicted that participants in the high-explore condition would make more exploratory choices.

Method

Participants. Thirty-one participants completed the experiment, which was conducted over Amazon Mechanical Turk (AMT) using the psiTurk framework (Gureckis et al., 2015). Participants were paid \$5.00 for their participation, with a performance-based bonus of up to \$3.00. All participants received the full \$3.00 bonus. Participants were pseudo-randomly counterbalanced across two conditions.

Design and procedure. *Consumption tasks.* Participants were informed that their job was to perform a monotonous slider task that would be split into 30-second “work periods,” but that they would be able to make choices throughout the experiment that would give them a chance to watch a YouTube video instead. The number of remaining work periods in the experiment was shown at the top of the screen, as was the number of seconds left in the current work period.

The slider task was based of a task previously used by Gill & Prowse (2012). In each period of the slider task, five horizontal sliders appeared on the screen (Figure 1a). Each started at a random setting between 0 and 100, with the slider's value shown to its right, and with a random horizontal offset so that the sliders were not aligned. The participant's task was to use the mouse to move each slider to “50” before the work period ended. When a participant released the mouse at the correct setting, the slider turned green to show it had been completed. To ensure that the task took close to the allotted 30 seconds, at the beginning of the task only the top slider was enabled, and the other four were grayed out. Additional sliders were enabled at five-second intervals, such that all five sliders were available after 20 seconds.

Before beginning the experiment, participants chose one of four videos available on YouTube: an episode of “Planet

Earth”, and episode of “The Great British Bakeoff”, and episode of “Mythbusters”, or an Ellen Degeneres comedy special. The video was embedded in the experiment with all user controls (such as skipping ahead) disabled (Figure 1b). When given access to the video, participants had to keep the browser window open and hold down the space bar for the video to play. This allowed us to ensure that participants maintained engagement with the content.

Participants completed a total of 70 work periods. For the first 10 work periods, participants simply clicked a button to begin the slider task. After these initial periods, participants gained access to six machines that could potentially complete the slider task for the participant, allowing the participant to watch their chosen video instead. However, the machines did not always function, and participants had to make a decision about how to set the machine before each use.

Decision-making task. Following the initial 10 periods, participants were shown a machine before each work period and, as shown in Figure 1c, had to select between two circular spinners with arrows at the top: the “current spinner” (exploit) and the “new spinner” (explore). The current spinner was split into five black and five gold wedges. If a participant chose the current spinner, it spun and, if the arrow landed on gold, the machine worked and the participant could watch the video. Initially, the current spinner's gold wedges were randomly set for each machine to cover between 1/3 and 2/3 of the spinner.

The new spinner initially showed a question mark. If a participant chose the new spinner, then a new spinner was created and appeared on the machine. In the high-explore condition, the new spinner's gold wedges could cover anywhere from 25% of the spinner to 100% of the spinner. In the low-explore condition, the new spinner's gold wedges could cover anywhere from 0% of the spinner to 75% of the spinner. The new spinner then spun and, if the arrow landed on gold, the machine worked.

Critically, if the new spinner had a greater gold area than the current spinner, the new spinner was “saved” and the current spinner was updated to the new spinner. This created an explore–exploit tradeoff in which choosing a new spinner carried immediate risk, but could carry long-term benefits by improving the current spinner from its initial value.

The machine ran immediately after the choice was made, and affected the next work period, as shown in Figure 2a and Figure 3a. It then “cooled off” for the following five periods, as choices were made with the other five machines.

Finally, in order to induce exploration throughout the entire experiment, the six machines would occasionally “reset” after they ran. When this occurred, the current spinner would be set to a new random value between 1/3 and 2/3 gold. Participants were informed that this would occur randomly on 1/6 of trials. In fact, the procedure was designed so exactly one of the six machines would reset on each cycle through the machines, and no machine would be reset on two consecutive uses.

Training, incentives, and post-experiment questions. Be-

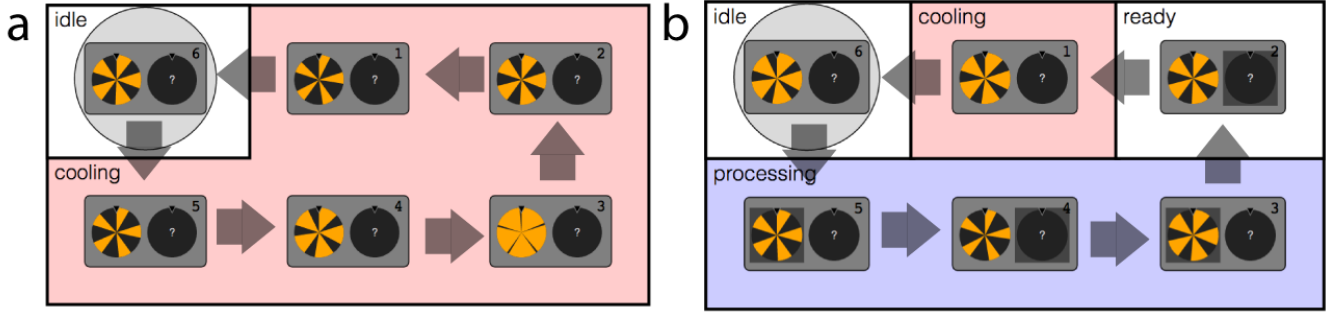


Figure 3: The machine display seen by participants. The display allowed participants track the value of each machine and the next time each machine would be ready to make a choice or produce an outcome. Gray arrows have been added to depict the counterclockwise movement of machines around the display after each work period. *a*: the display seen by participants in Experiment 1 and the immediate condition of Experiment 2. *b*: the display seen by participants in the delayed condition of Experiment 2.

fore beginning the full experiment, participants completed two practice phases. First, they performed several trials of practice using the machines, with the actual work periods removed. Then, they performed two work periods practicing the slider task and one work period practicing the video task. During the machine choices, participants had access to an “info” button at the bottom of the screen that provided reminders about the dynamics of the task.

Participants were given a performance-based bonus of \$3.00 for completing the consumption tasks accurately. If they missed fewer than 10% of sliders throughout the experiment and left the video paused less than 20% of the time, they were not penalized. However, if they missed more sliders or left the video paused for longer, they lost 10 cents from their bonus for each additional percentage of sliders missed or time with the video paused. The running percentage of sliders missed and video pause time was displayed at the top of the screen throughout the experiment.

Following the experiment, participants were asked to rate their enjoyment of the slider task and of the video-watching task on a 1 to 7 scale.

Results

Participants rated their enjoyment of the video at 5.77 out of 7, on average, and their enjoyment of the slider task at 2.87 out of 7, on average. This difference in enjoyment between the tasks was significant, $t(30) = 10.53$, $p < .001$.

To analyze participants’ trial-by-trial decision-making, we conducted a hierarchical Bayesian logistic regression using the Stan modeling language (Stan Development Team, 2015). This approach allowed us to estimate population-level effects of the current-spinner value and of condition, while also allowing for individual differences. The regression model included an intercept term as well as terms for the value of the current spinner, the participant’s condition, and a condition by value of current spinner interaction.

The results of the regression are shown in Figure 4. Because we observed few trials for each participant, and the current-spinner values observed for each participant depend on chance and the participant’s choices, we present

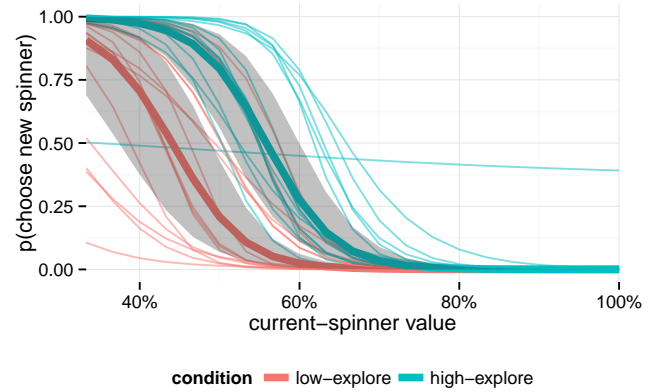


Figure 4: Model-based estimates of participants’ probability of choosing a new spinner for different values of the current spinner in Experiment 1. Thick lines and shaded regions indicate the mean and 95% posterior interval for the population-level parameters, while the thin lines indicate the mean posterior parameters for each of the 31 individual participants. Participants in the high-explore condition were more likely to explore at a given current-spinner value than those in the low-explore condition.

individual-level model fits rather than the raw choice data. We found that at the population level, participants were less likely to choose a new spinner when the value of the current spinner was higher, with posterior probability $p > .999$. We also found that participants in the high-explore were more likely to explore a new spinner for a given current spinner value than those in the low-explore condition, with posterior probability $p > .99$. We found no interaction between condition and the effect of current-spinner value.

These results show that participants’ behavior in this task is similar to behavior in other exploratory choice tasks. Participants explored more when the current exploitative option was poor, as well as when the possible payoffs of exploration were high. While this experiment did not address whether the consumption rewards produced a bias towards the present (a topic returned to in the discussion), it did show that participants enjoyed the video more than the slider task, and that they made choices to obtain more video time. In the follow-




Figure 5: Model-based estimates of participants’ probability of choosing a new spinner for different values of the current spinner in Experiment 2. Thick lines and shaded regions indicate the mean and 95% posterior interval for the population-level parameters, while the thin lines indicate the mean posterior parameters for each of the 40 individual participants. Participants in the delayed-outcome condition were no more likely to explore at a given current-spinner value than those in the immediate-outcome condition.

ing experiment, we introduce a delayed-outcome version of the task and predict that, if the consumption rewards induce a present bias, this version of the task will increase exploration.

Experiment 2

Method

Participants. Forty participants completed the experiment, which was conducted over AMT using the psiTurk framework. Participants were paid \$5.00 for their participation, with a performance-based bonus of up to \$3.00. All participants received the full \$3.00 bonus. Participants were pseudo-randomly counterbalanced across two conditions.

Design and procedure. Experiment 2 followed the design of Experiment 1, with the following changes. In both conditions, the possible winning proportion of a new spinner ranged from 0% to 100% of the spinner. In the immediate condition, the machine ran immediately after the choice was made, as in Experiment 1. In the delayed condition, each machine was presented to the participant four work periods before it was scheduled to run, and the participant made a choice at that time. The machine then had to “process” for four work periods, thus delaying the outcome by over 2 minutes (Figure 3b). The participant then returned to the machine to observe its outcome and either perform the slider task or watch the video. The machine then cooled off for a single period before being ready for another choice. The timeline for the delayed condition is shown in Figure 2b.

Results

As in Experiment 1, participants rated the video as more enjoyable on average (5.65 out of 7) than the slider task (3.13 out of 7), $t(39) = 9.26$, $p < .001$.

To analyze trial-by-trial decision-making, we conducted a hierarchical Bayesian logistic regression identical to the one described in the results of Experiment 1.

The results of the regression are shown in Figure 5. As in Experiment 1, participants were less likely to choose a new spinner when the current spinner has a high value, $p > .999$. However, in this experiment we found no evidence of an effect of condition, with a posterior probability of $p = .558$ that the effect of condition was greater than 0. This means that participants were no more likely to explore a new spinner when there was a temporal delay imposed between their choices and the received outcomes. However, there may have been a small interaction between current spinner value and condition, such that participants in the delayed condition were less sensitive to the value of the spinner when making their choices. This might indicate that the delayed condition was confusing to some participants, as a few individuals (as seen in Figure 5) changed their behavior very little across current-spinner values. The posterior probability that this interaction was above 0 is $p = .934$.

Discussion

In this paper we presented a novel task using immediately-consumed rewards to study present bias in the context of exploratory choice. We predicted that adding a temporal delay between choices and outcomes would increase exploration by reducing the weight of the current choice’s outcome (which is maximized through exploitation) relative to future choices’ outcomes (which are increased through exploration). In Experiment 1 we showed that participants appeared to understand the task and respond reasonably to maximize their time in the video-watching task. They were more likely to explore when the current exploitable option (current spinner) had low value, and changed their degree of exploration based on the range of possible outcomes from choosing a new spinner. However, in contrast to our prediction, in Experiment 2 we found that participants did not explore any more when there was a temporal delay between choices and outcomes.

One explanation for our results is that people are in fact less affected by present bias in the realm of exploratory choice than they are in other inter-temporal choices. This might be the case if exploratory choice is driven less by an explicit present–future tradeoff and more by intrinsic “exploration bonuses” (Kakade & Dayan, 2002) or motivational states such as curiosity (Berlyne, 1966). If exploration is driven by these mechanisms, then the reward of exploration might be exploration itself rather than its later consequences, breaking the intertemporal tradeoff between exploration and exploitation.

However, given the early stage of the present work, we are not ready to make strong claims about the role of present bias in exploratory choice. First, while participants enjoyed the video more than the slider task, we need to perform a follow-up study to ensure that this difference in enjoyment was strong enough to create a bias towards immediate video-

watching in the time scale used in our experiments. To test this, we might use a simple experiment in which two buttons produce video-watching time at for different lengths (or with different probabilities) and at different delays. In this setting, we would predict a bias towards the button that produced earlier video-watching time, even if the video time was shorter or less frequent. Navarick (1998) used a similar paradigm to demonstrate present bias in watching cartoons, though the bias might be reduced in participants on AMT due to the greater availability of environmental distractors during non-video time.

We also may need to improve our decision-making task. As discussed, some participants may have found the delayed condition confusing, and this problem might be fixed through clearer instructions or interface. Additionally, the future encounters with each machine may not have been sufficiently salient, reducing the perceived need for exploration in both conditions.

Acknowledgments This research was supported by NSF grant BCS-1255538 and a John S. McDonnell Foundation Scholar Award to TMG.

References

- Abramson, L. Y., Seligman, M. E., & Teasdale, J. D. (1978). Learned helplessness in humans: critique and reformulation. *Journal of abnormal psychology*, 87(1), 49–74. doi: 10.1037/0021-843X.87.1.49
- Bartels, D. M., & Urminsky, O. (2015). To Know and to Care: How Awareness and Valuation of the Future Jointly Shape Consumer Spending. *Journal of Consumer Research*, 41(6), 1469–1485. doi: 10.1086/680670
- Berlyne, D. E. (1966). Curiosity and Exploration. *Science*.
- Brown, A. L., Chua, Z. E., & Camerer, C. F. (2009). Learning and visceral temptations in dynamic saving experiments. *The Quarterly Journal of Economics*(February), 197–231.
- Denrell, J. (2005). Why most people disapprove of me: experience sampling in impression formation. *Psychological review*, 112(4), 951–978.
- Fang, C., Lee, J., & Schilling, M. a. (2010). Balancing Exploration and Exploitation Through Structural Design: The Isolation of Subgroups and Organizational Learning. *Organization Science*, 21(3), 625–642. doi: 10.1287/orsc.1090.0468
- Gill, D., & Prowse, V. (2012). A Structural Analysis of Disappointment Aversion in a Real Effort Competition. *The American Economic Review*, 102(1), 469–503.
- Gopher, D., Weil, M., & Siegel, D. (1989). Practice under changing priorities: An approach to the training of complex skills. *Acta Psychologica*, 71(1-3), 147–177. doi: 10.1016/0001-6918(89)90007-3
- Guez, A., Silver, D., & Dayan, P. (2013). Scalable and efficient bayes-adaptive reinforcement learning based on Monte-Carlo tree search. *Journal of Artificial Intelligence Research*, 48, 841–883. doi: 10.1613/jair.4117
- Gureckis, T. M., Martin, J., McDonnell, J., Rich, A. S., Markant, D., Coenen, A., ... Chan, P. (2015). psiTurk: An open-source framework for conducting replicable behavioral experiments online. *Behavior Research Methods*. doi: 10.3758/s13428-015-0642-8
- Hertwig, R., Barron, G., Weber, E. U., & Erev, I. (2004). Decisions from experience and the effect of rare events in risky choice. *Psychological Science*, 15(8), 534–539.
- Huys, Q. J. M., & Dayan, P. (2009). A Bayesian formulation of behavioral control. *Cognition*, 113(3), 314–328. doi: 10.1016/j.cognition.2009.01.008
- Juni, M. Z., Gureckis, T. M., & Maloney, L. T. (2016). Information Sampling Behavior With Explicit Sampling Costs. *Decision*.
- Kakade, S., & Dayan, P. (2002). Dopamine: generalization and bonuses. *Neural Networks*, 15(4), 549–559.
- Kirby, K. N., & Herrnstein, R. (1995). Preference Reversals Due To Myopic Discounting of Delayed Reward. *Psychological Science*, 6(2), 83–89. doi: 10.1111/j.1467-9280.1995.tb00311.x
- Levinthal, D. a., & March, J. G. (1993). The myopia of learning. *Strategic Management Journal*, 14, 95–112. doi: 10.1002/smj.4250141009
- March, J. G. (1991). Exploration and Exploitation in Organizational Learning. *Organization Science*, 2(1), 71–87.
- Mehlhorn, K., Newell, B. R., Todd, P. M., Lee, M. D., Morgan, K., Braithwaite, V. A., ... Fiedler, K. (2015). Unpacking the Exploration Exploitation Tradeoff : A Synthesis of Human and Animal Literatures. *Decision*, 2(3), 191–215.
- Navarick, D. (1998). Impulsive choice in adults: How consistent are individual differences? *The Psychological Record*, 48, 665–674.
- Navarro, D. J., Newell, B. R., & Schulze, C. (2016). Learning and choosing in an uncertain world : An investigation of the explore-exploit dilemma in static and dynamic environments. *Cognitive Psychology*, 85, 43–77. doi: 10.1016/j.cogpsych.2016.01.001
- Rich, A. S., & Gureckis, T. M. (2017). Exploratory choice reflects the future value of information. *Decision*, (in press).
- Sang, K., Todd, P. M., & Goldstone, R. L. (2011). Learning near-optimal search in a minimal explore/exploit task. *Proceedings of the Thirty-third Annual Conference of the Cognitive Science Society*, 2800–2805.
- Shook, N. J., & Fazio, R. H. (2008). Interracial roommate relationships: An experimental field test of the contact hypothesis: Research article. *Psychological Science*, 19(7), 717–723. doi: 10.1111/j.1467-9280.2008.02147.x
- Solnick, J. V., Kannenberg, C. H., Eckerman, D. A., & Waller, M. B. (1980). An experimental analysis of impulsivity and impulse control in humans. *Learning and Motivation*, 11(1), 61–77. doi: 10.1016/0023-9690(80)90021-1
- Stan Development Team. (2015). *Stan: A C++ Library for Probability and Sampling, Version 2.7*. Retrieved from <http://mc-stan.org/>
- Sutton, R. S., & Barto, A. G. (1998). *Reinforcement learning: An introduction*. Cambridge Univ Press.
- Teodorescu, K., & Erev, I. (2014a). Learned Helplessness and Learned Prevalence: Exploring the Causal Relations Among Perceived Controllability, Reward Prevalence, and Exploration. *Psychological Science*, 25(10), 1861–1869. doi: 10.1177/0956797614543022
- Teodorescu, K., & Erev, I. (2014b). On the Decision to Explore New Alternatives: The Coexistence of Under- and Over-exploration. *Journal of Behavioral Decision Making*, 27, . doi: 10.1002/bdm
- Tversky, A., & Edwards, W. (1966). Information versus reward in binary choices. *Journal of Experimental Psychology*, 71(5), 680.
- Wilson, R. C., Geana, A., White, J. M., Ludvig, E. A., & Cohen, J. D. (2014). Humans Use Directed and Random Exploration to Solve the Explore Exploit Dilemma. *Journal of Experimental Psychology: General*.
- Yechiam, E., Erev, I., & Gopher, D. (2001). On the potential value and limitations of emphasis change and other exploration-enhancing training methods. *Journal of experimental psychology. Applied*, 7(4), 277–285. doi: 10.1037/1076-898X.7.4.277
- Zwick, R., Rapoport, A., Lo, A. K. C., & Muthukrishnan, a. V. (2003). Consumer Sequential Search: Not Enough or Too Much? *Marketing Science*, 22(4), 503–519. doi: 10.1287/mksc.22.4.503.24909